

Provider	Model	Release Date	Knowledge Cutoff	Cost (Input / Output per 1M tokens)	Context Window	Public Model	When to Use (Key Strength)
Anthropic	Claude Opus 4.7	Apr 16, 2026 (Venturebeat)	~Late 2025 (est.)	\$5 / \$25 (DataCamp)	1M input / 128K output (DataCamp)	Yes	Top-tier reasoning, scientific + enterprise workflows
OpenAI	GPT-5.4	~Mar 2026 (est., relative timing) (DocsBot AI)	~Late 2025	\$2.5 / \$15 (DocsBot AI)	~1M+ (DocsBot AI)	Yes	Balanced reasoning + long-context + production apps
Anthropic	Claude Haiku 4.5	Oct 2025 (Artificial Analysis)	~Mid 2025	~\$1 / \$5 (DocsBot AI)	200K (Artificial Analysis)	Yes	Fast, cheap, real-time apps (chatbots, agents)
xAI	Grok Code Fast 1	~2026 (limited public info)	—	—	—	Partial	Low-latency coding + dev assistants (fast inference)
OpenAI	GPT-5.2-Codex	~Early 2026 (inferred from GPT-5.2 cycle) (AceCloud)	~2025	~\$2–\$10 (varies by tier) (AceCloud)	~400K	Yes	Code generation + agentic workflows
Google DeepMind	Gemini 3.1 Pro	Feb 19, 2026 (EvoLink)	~2025	\$2 / \$12 (DataCamp)	1M (DataCamp)	Yes (preview)	Best for multimodal + large-scale reasoning + agents
OpenAI	GPT-5.3-Codex	~Feb–Mar 2026 (inferred) (Digital Applied)	~2025	Not fully public	~400K	Yes	Best for terminal-heavy coding +

Provider	Model	Release Date	Knowledge Cutoff	Cost (Input / Output per 1M tokens)	Context Window	Public Model	When to Use (Key Strength)
							automation agents
Anthropic	Claude Opus 4.6	Feb 5, 2026 (Anthropic)	~Late 2025	~\$5 / \$25 (similar tier)	1M (Medium)	Yes	Agentic coding + long-running workflows
OpenAI	GPT-5 Codex (base)	Sep 2025 (Artificial Analysis)	~2025	—	400K (Artificial Analysis)	Yes	General coding + dev tooling baseline
Anthropic	Claude Sonnet 4.5 (<i>contextual reference</i>)	2025	~2025	—	—	Yes	Balanced reasoning + coding (mid-tier)

Provider	Model	Release Date	Knowledge Cutoff	Cost (Input / Output per 1M tokens)	Context Window	Public Model	When to Use (Key Strength)
OpenAI	GPT-5.5	Apr 2026	~Nov 2025 (est.)	Not fully public yet	~400K (est.)	Yes	Best overall reasoning + coding + agent tasks
OpenAI	GPT-5.2	Dec 2025	~Oct 2025	\$1.75 / \$14	400K	Yes	Advanced reasoning, math-heavy tasks
OpenAI	o3 (reasoning model)	Apr 2025	~2025	\$10 / \$40	200K	Yes	Deep reasoning, step-by-step thinking (expensive)
OpenAI	GPT-4.1	Apr 2025	~2024	\$2 / \$8	1M	Yes	Long-context apps (documents, RAG)

Provider	Model	Release Date	Knowledge Cutoff	Cost (Input / Output per 1M tokens)	Context Window	Public Model	When to Use (Key Strength)
Anthropic	Claude Opus 4.6	Feb 2026	~Late 2025	\$5 / \$25	200K–1M	Yes	Enterprise reasoning + long workflows
Anthropic	Claude Sonnet 4.6	2026	~2025	\$3 / \$15	200K–1M	Yes	Best cost vs performance (balanced tasks)
Google	Gemini 3 Pro	Nov 2025	~2025	~\$2 / \$12	1M+	Yes	Multimodal + long context + research tasks
Google	Gemini 2.5 Pro	2025	~2024	\$1.25 / \$10	1M	Yes	High-quality + cheaper than frontier models
Meta	Llama 4 Scout	2026	Open weights	Free (self-host)	up to 10M	Yes (open)	Ultra-long context + on-prem deployment
Mistral	Mistral Large	2025	~2024	\$2 / \$6	128K	Yes	Strong multilingual + cheaper enterprise option